

Architectural anti-patterns when delivering a software ecosystem with Kubernetes

Contents

Architectural anti-patterns when delivering a software ecosystem with Kubernetes	1
1. Setup	2
2. Saying hello – Kubernetes components, orchestration and more.....	6
3. Statefulness – the first evil of all microservices architectures	11
4. Privacy – mind your secrets	12
5. The undo rollout – or how to deal with the dependencies at undo rollout	14
6. The actuator – check the healthiness of your container	20
7. An unfortunate choice - having multiple containers in the same pod	23
8. Limits and auto-scaling	24

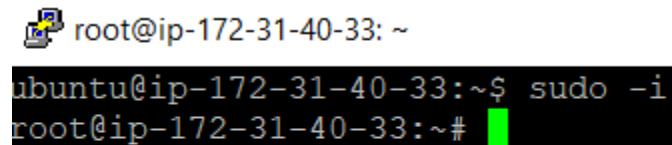
1. Setup

In order to be able to solve all the examples we must first set up our environments.

The first step is to create your own EC2 virtual machine in the space provided by the trainer. You can as well choose to use your own computer if you have the needed privileges. However, be aware that the commands might be different in some situations if you use your own computer – mostly depending on the OS that you have installed. The EC2 will have installed an Ubuntu server 16.04, so the commands in this document are focused on this OS version.

After first accessing the EC2 virtual machine, let's make sure we have root permissions. You can do this by running:

sudo -i

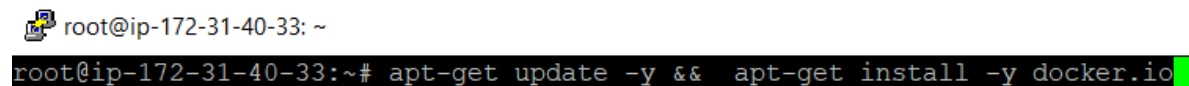


```
root@ip-172-31-40-33: ~  
ubuntu@ip-172-31-40-33:~$ sudo -i  
root@ip-172-31-40-33:~#
```

Next step is to install Docker. To achieve this, please run the following command:

apt-get update -y

apt-get install -y docker.io



```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# apt-get update -y && apt-get install -y docker.io
```

If there is no error message, please check your installation using the *docker version* command as shown below:

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# docker version  
Client:  
Version:      17.03.2-ce  
API version:  1.27  
Go version:   go1.6.2  
Git commit:   f5ec1e2  
Built:        Thu Jul  5 23:07:48 2018  
OS/Arch:      linux/amd64  
  
Server:  
Version:      17.03.2-ce  
API version:  1.27 (minimum version 1.12)  
Go version:   go1.6.2  
Git commit:   f5ec1e2  
Built:        Thu Jul  5 23:07:48 2018  
OS/Arch:      linux/amd64  
Experimental: false  
root@ip-172-31-40-33:~#
```

If everything is fine, we should now install Minikube:

```
curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64 &&  
chmod +x minikube && sudo mv minikube /usr/local/bin/
```

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# curl -Lo minikube https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64 && chmod +x minikube && sudo mv minikube /usr/local/bin/
```

In order to use our Kubernetes cluster on Minikube, we need the Kubernetes command line tool installed. This is kubectl. To install kubectl, run this command:

```
curl -LO https://storage.googleapis.com/kubernetes-release/release/v1.10.0/bin/linux/amd64/kubectl  
&& chmod +x ./kubectl && sudo mv ./kubectl /usr/local/bin/kubectl
```

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# curl -Lo kubectl https://storage.googleapis.com/kubernetes-release/release/v1.8.0/bin/linux/amd64/kubectl && chmod +x kubectl && sudo mv kubectl /usr/local/bin/
```

Next, we should install git, maven and the JDK:

```
apt-get install -y git
```

```
apt-get install -y maven
```

```
apt-get install -y default-jdk
```

Check that they have been installed correctly using:

```
java -version
```

```
mvn -version
```

```
git --version
```

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# java -version  
openjdk version "1.8.0_191"  
OpenJDK Runtime Environment (build 1.8.0_191-8u191-b12-0ubuntu0.16.04.1-b12)  
OpenJDK 64-Bit Server VM (build 25.191-b12, mixed mode)  
root@ip-172-31-40-33:~# mvn -version  
Apache Maven 3.3.9  
Maven home: /usr/share/maven  
Java version: 1.8.0_191, vendor: Oracle Corporation  
Java home: /usr/lib/jvm/java-8-openjdk-amd64/jre  
Default locale: en_US, platform encoding: UTF-8  
OS name: "linux", version: "4.4.0-1072-aws", arch: "amd64", family: "unix"  
root@ip-172-31-40-33:~# git --version  
git version 2.7.4  
root@ip-172-31-40-33:~# █
```

At this moment we could also login to the Docker account where we will push our images.

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# docker login  
Login with your Docker ID to push and pull images from Docker Hub.  
m to create one.  
Username: █
```

Start minikube with this command.

```
minikube start --memory=8192 --cpus=4 --vm-driver=none
```

Then check that the service works fine by giving some commands to check the current deployments, pods and services:

```
root@ip-172-31-40-33: ~  
root@ip-172-31-40-33:~# kubectl get deployments  
No resources found.  
root@ip-172-31-40-33:~# kubectl get pods  
No resources found.  
root@ip-172-31-40-33:~# kubectl get services  
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE  
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP    5m  
root@ip-172-31-40-33:~#
```

During the tutorial we will need metrics, especially when dealing with horizontal auto scaling. First step is to install the metrics server on Kubernetes. In order to do this, you have to enable metrics-server, clone the following git repository and apply the files in the deploy/1.8+ folder:

```
minikube addons enable metrics-server
```

```
git clone https://github.com/kubernetes-incubator/metrics-server.git
```

```
kubectl apply -f deploy/1.8+
```

2. Saying hello – Kubernetes components, orchestration and more...

First, we can check all *docker images*.

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
gcr.io/knative-releases/github.com/knative/eventing/pkg/buses/stub/dispatcher	<none>	c4db305beca6	4 weeks ago	55.9 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/webhook	<none>	62ae21262386	4 weeks ago	56.2 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/controller	<none>	d93a517e25ef	4 weeks ago	61 MB
openjdk	8-jdk-alpine	97bc1352afde	5 weeks ago	103 MB
k8s.gcr.io/kube-proxy	v1.12.2	15e9da1ca195	5 weeks ago	96.5 MB
k8s.gcr.io/kube-apiserver	v1.12.2	51a9c329b7c5	5 weeks ago	194 MB
k8s.gcr.io/kube-controller-manager	v1.12.2	15548c720a70	5 weeks ago	164 MB
k8s.gcr.io/kube-scheduler	v1.12.2	d6d57c76136c	5 weeks ago	58.3 MB
k8s.gcr.io/etcd	3.2.24	3cab8e1b9802	2 months ago	220 MB
istio/galley	1.0.2	b8cfc0e19a91	2 months ago	65.8 MB
istio/citadel	1.0.2	ca4050c9fed3	2 months ago	50.7 MB
istio/mixer	1.0.2	d559bdcd7a88	2 months ago	64.5 MB
istio/sidecar_injector	1.0.2	77e6870301bb	2 months ago	45.3 MB
istio/proxy_init	1.0.2	4cd353237d97	2 months ago	119 MB
istio/proxyv2	1.0.2	50d4ec2a16fd	2 months ago	371 MB
istio/pilot	1.0.2	3be7ec27d893	2 months ago	308 MB
k8s.gcr.io/coredns	1.2.2	367cdc8433a4	3 months ago	39.2 MB
k8s.gcr.io/kubernetes-dashboard-amd64	v1.10.0	0dab2435c100	3 months ago	122 MB
k8s.gcr.io/kube-addon-manager	v8.6	9c16409588eb	9 months ago	78.4 MB
prom/statsd-exporter	v0.6.0	304735eab4e4	10 months ago	14.1 MB
k8s.gcr.io/pause	3.1	da86e6ba6ca1	11 months ago	742 kB
gcr.io/k8s-minikube/storage-provisioner	v1.8.1	4689081edb10	12 months ago	80.8 MB
quay.io/coreos/hyperkube	v1.7.6_coreos.0	2faf6f7a322f	14 months ago	699 MB
gcr.io/knative-releases/github.com/knative/build/cmd/controller	<none>	bed10b848745	48 years ago	51.3 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/activator	<none>	ce1d33e5dfe9	48 years ago	53.2 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/autoscaler	<none>	98f79403ef44	48 years ago	54.7 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/webhook	<none>	5bd6d6d43479	48 years ago	51.8 MB
gcr.io/knative-releases/github.com/knative/build/cmd/webhook	<none>	2c475185743c	48 years ago	49.4 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/controller	<none>	2c57b73e7aac	48 years ago	56.8 MB

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1#
```

And then we can also check the deployments and the services again. On the default namespace, there should be no deployment yet and we should be able to see only the kubernetes service itself.

kubectl get pods

kubectl get deployments

kubectl get services

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1# kubectl get pods
```

No resources found.

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1# kubectl get deployments
```

No resources found.

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1# kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	52m

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1#
```

Now, let's change the directory to the /var and clone the git repository provided:

```

root@ip-172-31-40-33: /var
root@ip-172-31-40-33:/var# git clone https://github.com/lspil/saconf2019.git
Cloning into 'saconf2019'...
remote: Enumerating objects: 193, done.
remote: Counting objects: 100% (193/193), done.
remote: Compressing objects: 100% (103/103), done.
remote: Total 193 (delta 63), reused 167 (delta 37), pack-reused 0
Receiving objects: 100% (193/193), 59.25 KiB | 0 bytes/s, done.
Resolving deltas: 100% (63/63), done.
Checking connectivity... done.
root@ip-172-31-40-33:/var# ls
backups  cache  crash  lib  local  lock  log  mail  opt  run  saconf2019  snap  spool  tmp
root@ip-172-31-40-33:/var#

```

You can change dir now to saconf2019 where you should see all the examples:

```

root@ip-172-31-40-33: /var/saconf2019
root@ip-172-31-40-33:/var# cd saconf2019/
root@ip-172-31-40-33:/var/saconf2019# ls
saconf2019-e1  saconf2019-e2  saconf2019-e3
root@ip-172-31-40-33:/var/saconf2019#

```

In the folder of the first example you can find a Dockerfile. The content of the Dockerfile presents a very simple setup of the image that we wish to create. The image starts from the initial layer of the openjdk 8 alpine and adds a jar file provided through a build argument. Then it starts the application.

```

ubuntu@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
ubuntu@ip-172-31-40-33:/var/saconf2019$ cd saconf2019-e1
ubuntu@ip-172-31-40-33:/var/saconf2019/saconf2019-e1$ ls
Dockerfile  kube  mvnw  mvnw.cmd  pom.xml  src
ubuntu@ip-172-31-40-33:/var/saconf2019/saconf2019-e1$ cat Dockerfile
FROM openjdk:8-jdk-alpine
VOLUME /tmp
ARG JAR_FILE
COPY ${JAR_FILE} app.jar
ENTRYPOINT ["java","-Djava.security.egd=file:/dev/./urandom","-jar","/app.jar"]
ubuntu@ip-172-31-40-33:/var/saconf2019/saconf2019-e1$

```

The kube folder contains the yml files with the description of the deployments, services, secrets etc.

```

ubuntu@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
ubuntu@ip-172-31-40-33:/var/saconf2019/saconf2019-e1$ ls kube
deployment.yml  service.yml
ubuntu@ip-172-31-40-33:/var/saconf2019/saconf2019-e1$

```

Next step needed for running an application in Kubernetes is having the Docker image that will be used to create the running containers. Providing a Docker image, Kubernetes will create the pods according to the deployment yml file.

To create the Docker image, we will use the docker file. Running the docker build command, the only thing we need to provide as a parameter is the value of the JAR_FILE build argument. To obtain the jar file, we simply have to compile the application using Maven:

mvn clean install

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# ls
Dockerfile kube mvnw mvnw.cmd pom.xml src target
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# ls target/
classes generated-test-sources maven-status saconf2019-e1-0.0.1-SNAPSHOT.jar.original test-classes
generated-sources maven-archiver saconf2019-e1-0.0.1-SNAPSHOT.jar surefire-reports
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1#
```

After a successful build the target folder containing the fat Spring boot jar should appear.

We can use it to create the Docker image

docker build . --build-arg=JAR_FILE=target/saconf2019-e1-0.0.1-SNAPSHOT.jar

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# ls target/
classes generated-test-sources maven-status saconf2019-e1-0.0.1-SNAPSHOT.jar.original test-classes
generated-sources maven-archiver saconf2019-e1-0.0.1-SNAPSHOT.jar surefire-reports
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# docker build . --build-arg=JAR_FILE=target/saconf2019-e1-0.0.1-SNAPSHOT.jar
```

Then we should be able to see it within the *docker images*


```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# docker images
REPOSITORY                                TAG                                IMAGE ID                                CREATED                                SIZE
<none>                                    <none>                            e930db8f4ed5                          53 seconds ago                       119 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/webhook                <none>                            62ae21262386                          4 weeks ago                           56.2 MB
gcr.io/knative-releases/github.com/knative/eventing/pkg/buses/stub/dispatcher <none>                            c4db305beca6                          4 weeks ago                           55.9 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/controller              <none>                            d93a517e25ef                          4 weeks ago                           61 MB
openjdk                                   8-jdk-alpine                     97bc1352afde                          5 weeks ago                           103 MB
k8s.gcr.io/kube-proxy                    v1.12.2                          15e9dalc195                           5 weeks ago                           96.5 MB
k8s.gcr.io/kube-apiserver                 v1.12.2                          51a9c329b7c5                          5 weeks ago                           194 MB
k8s.gcr.io/kube-controller-manager        v1.12.2                          15548c720a70                          5 weeks ago                           164 MB
k8s.gcr.io/kube-scheduler                 v1.12.2                          d6d57c76136c                          5 weeks ago                           58.3 MB
k8s.gcr.io/etcd                           3.2.24                           3cab8e1b9802                          2 months ago                          220 MB
istio/galley                             1.0.2                             b8cfc0e19a91                          2 months ago                          65.8 MB
istio/citadel                            1.0.2                             ca4050c9fed3                          2 months ago                          50.7 MB
istio/mixer                              1.0.2                             d559bdcd7a88                          2 months ago                          64.5 MB
istio/sidecar_injector                   1.0.2                             77e6870301bb                          2 months ago                          45.3 MB
istio/proxy_init                         1.0.2                             4cd353237d97                          2 months ago                          119 MB
istio/proxyv2                            1.0.2                             50d4ec2a16fd                          2 months ago                          371 MB
istio/pilot                              1.0.2                             3be7ec27d893                          2 months ago                          308 MB
k8s.gcr.io/coredns                       1.2.2                             367cdc8433a4                          3 months ago                          39.2 MB
k8s.gcr.io/kubernetes-dashboard-amd64    v1.10.0                           0dab2435c100                          3 months ago                          122 MB
k8s.gcr.io/kube-addon-manager             v8.6                              9c16409588eb                          9 months ago                          78.4 MB
prom/statsd-exporter                     v0.6.0                            304735eab4e4                          10 months ago                         14.1 MB
k8s.gcr.io/pause                          3.1                               da86e6ba6ca1                          11 months ago                         742 kB
gcr.io/k8s-minikube/storage-provisioner   v1.8.1                            4689081edb10                          12 months ago                         80.8 MB
quay.io/coreos/hyperkube                 v1.7.6_coreos.0                  2faf6f7a322f                          14 months ago                         699 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/controller              <none>                            2c57b73e7aac                          48 years ago                           56.8 MB
gcr.io/knative-releases/github.com/knative/build/cmd/webhook                  <none>                            2c475185743c                          48 years ago                           49.4 MB
gcr.io/knative-releases/github.com/knative/build/cmd/controller                <none>                            bed10b8484745                          48 years ago                           51.3 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/webhook                 <none>                            5bd6d6d43479                          48 years ago                           51.8 MB
gcr.io/knative-releases/github.com/knative/serving/cmd/activator               <none>                            ce1d33e5dfe9                          48 years ago                           53.2 MB
gcr.io/knative-releases/github.com/knative/autoscaler                         <none>                            98f79403ef44                          48 years ago                           54.7 MB
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1#

```

The last image, created a couple of seconds ago, is currently untagged and it is the one created by the previous command. Now it can be tagged and pushed to the repository.

`docker tag e930db8f4ed5 laurentiuspilca/saconf2019-e1:v1`

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# docker images
REPOSITORY                                TAG                                IMAGE ID                                CREATED                                SIZE
laurentiuspilca/saconf2019-e1            v1                                e930db8f4ed5                          6 minutes ago                         119 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/controller              <none>                            d93a517e25ef                          4 weeks ago                           61 MB
gcr.io/knative-releases/github.com/knative/eventing/cmd/webhook                <none>                            62ae21262386                          4 weeks ago                           56.2 MB
gcr.io/knative-releases/github.com/knative/eventing/pkg/buses/stub/dispatcher <none>                            c4db305beca6                          4 weeks ago                           55.9 MB
openjdk                                   8-jdk-alpine                     97bc1352afde                          5 weeks ago                           103 MB
k8s.gcr.io/kube-proxy                    v1.12.2                          15e9dalc195                           5 weeks ago                           96.5 MB
k8s.gcr.io/kube-apiserver                 v1.12.2                          51a9c329b7c5                          5 weeks ago                           194 MB
k8s.gcr.io/kube-controller-manager        v1.12.2                          15548c720a70                          5 weeks ago                           164 MB
k8s.gcr.io/kube-scheduler                 v1.12.2                          d6d57c76136c                          5 weeks ago                           58.3 MB
k8s.gcr.io/etcd                           3.2.24                           3cab8e1b9802                          2 months ago                          220 MB
istio/galley                             1.0.2                             b8cfc0e19a91                          2 months ago                          65.8 MB

```

`docker push laurentiuspilca/saconf2019-e1:v1`

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# docker push laurentiuspilca/saconf2019-e1:v1
The push refers to a repository [docker.io/laurentiuspilca/saconf2019-e1]
cfd5376efc00: Pushing [=====>] 16.24 MB
ed6f0bd39121: Layer already exists
0c3170905795: Layer already exists
df64d3292fd6: Layer already exists

```

`kubectl apply -f target/deployment.yml`

`kubectl apply -f target/service.yml`

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# ls
Dockerfile kube mvnw mvnw.cmd pom.xml src target
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl apply -f kube/deployment.yml
deployment "hello-deployment" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl apply -f kube/service.yml
service "hello-service" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# █

```

Let's check again the deployment, pods, and services.

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
hello-deployment-5c8f864485-hvqrz  1/1     Running   0           1m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get deployments
NAME                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello-deployment    1          1          1             1           1m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get services
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
hello-service       LoadBalancer  10.105.30.56 <pending>     8080:31108/TCP   1m
kubernetes           ClusterIP      10.96.0.1    <none>        443/TCP          1h
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# █

```

curl http://localhost:port/hello

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e1
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
hello-deployment-5c8f864485-hvqrz  1/1     Running   0           1m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get deployments
NAME                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
hello-deployment    1          1          1             1           1m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# kubectl get services
NAME                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
hello-service       LoadBalancer  10.105.30.56 <pending>     8080:31108/TCP   1m
kubernetes           ClusterIP      10.96.0.1    <none>        443/TCP          1h
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# curl http://localhost:31108/hello
Helloroot@ip-172-31-40-33:/var/saconf2019/saconf2019-e1# █

```

3. Statefulness – the first evil of all microservices architectures

To run the second example simply assume the Docker images are there. Just apply the deployment and the service.

```
kubectl apply -f target/deployment.yml
```

```
kubectl apply -f target/service.yml
```

After running the commands you will see 10 replicas of the same pod starting. You can play with the number of replicas to make it smaller or bigger by changing the replicas parameter in the deployment.yml file.

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e2
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e2# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-deployment-5c8f864485-hvqrz	1/1	Running	0	30m
stateful-app-deployment-689845fbc6-46rzw	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-88q92	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-8cglv	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-8vp84	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-bn496	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-jknxf	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-lnjcr	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-t5rbl	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-tf9hm	1/1	Running	0	16m
stateful-app-deployment-689845fbc6-wvg9f	1/1	Running	0	16m

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e2#
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e2
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e2# kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-service	LoadBalancer	10.105.30.56	<pending>	8080:31108/TCP	23m
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	2h
stateful-app	LoadBalancer	10.102.166.130	<pending>	8080:32703/TCP	9m

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e2#
```

The second example has two endpoints. One of them is setting a name. The other is returning hello to that name. The name is stored in the state of the container. Let's see what happens when we set a name and then say hello multiple times using more than one replica.

```
curl -XPOST http://localhost:32703/name/John
```

```
curl http://localhost:32703/hello
```

The observation should state the conclusion that keeping state on the application will result in keeping state only on one pod. But, as we want an architecture horizontally scalable this becomes a problem. When we scale our pods to multiple instances, requests are now spread over the replicas. This means that two consecutive requests might not reach the same pod.

4. Privacy – mind your secrets

When we deploy the system using an orchestration tool in a cloud it is always important to know where to keep the sensitive data like users, password or encryption keys used by the deployed applications.

In Kubernetes we keep such data in secrets. In this case, to be easier to use, the secret is also defined as a yml file. But you would not do this in a real case scenario and of course neither you should store them in git.

In the following example you can see one way in which the application can read the secret and use it.

```
cat secret.yml
```

Observe in the yml file that the value is base 64 encoded. Copy the value and decode it.

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e3/kube
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3/kube# cat secret.yml
apiVersion: v1
kind: Secret
metadata:
  name: saconf2019-ex3-secret
type: Opaque
data:
  my.secret.name: Sm9obg==root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3/kube#
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3/kube#
```

```
kubectl apply -f kube/secret.yml
```

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e3
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3# kubectl apply -f kube/secret.yml
secret "saconf2019-ex3-secret" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3#

```

Once the secret is created, apply the yml configuration for deployment and service.

```
kubectl apply -f kube/service.yml
```

```
kubectl apply -f kube/deployment.yml
```

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e3
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3# kubectl apply -f kube/service.yml
service "secret-app-service" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3# kubectl apply -f kube/deployment.yml
deployment "secret-app-deployment" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3#

```

```
kubectl get services
```

After calling the endpoint in the application, you can see that the value of the secret is used.

```
curl http://localhost:<port>
```

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e3
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3# kubectl get services
NAME                TYPE          CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes           ClusterIP     10.96.0.1        <none>            443/TCP          6d
secret-app-service   LoadBalancer 10.107.209.45    <pending>         8080:31662/TCP   2m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e3# curl http://localhost:31662
My secret name is: Johnroot@ip-172-31-40-33:/var/saconf2019/saconf2019-e3#

```

5. The undo rollout – or how to deal with the dependencies at undo rollout

When thinking about the architecture of the system, it is always important to understand what happens with the dependencies at an update. An orchestration tool like Kubernetes allows us to rollout undo the deployment. Because orchestration is “playing” with images, we store the last snapshot of our application so it is fairly easy to undo to an earlier point. However, depending on the architecture of the system we must understand if rollout and undo of the java code is just enough.

In the below example you find a situation in which the rollout undo is not enough.

We start by deploying a mysql server.

```
kubectl apply -f kube/mysql-pv.yml
```

```
kubectl apply -f kube/mysql-deployment.yml
```

```
kubectl apply -f kube/mysql-service.yml
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4# kubectl apply -f kube/mysql-pv.yml
persistentvolume "mysql-pv-volume" configured
persistentvolumeclaim "mysql-pv-claim" configured
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4# kubectl apply -f kube/mysql-deployment.yml
deployment "mysql" created
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4# kubectl apply -f kube/mysql-service.yml
service "mysql" created
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4#
```

```
kubectl get pods
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
mysql-6b698dfcbb-q8vwm             1/1     Running   0           42s
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4# kubectl get services
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kubernetes  ClusterIP   10.96.0.1    <none>        443/TCP          7d
mysql      LoadBalancer 10.105.65.52 <pending>     3306:30001/TCP   43s
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4#
```

Let's now use a client to connect to the server.

```
kubectl run -it --rm --image=mysql:5.6 --restart=Never mysql-client -- mysql -h mysql -ppassword
```

So we can create a database which we connect to.

create database saconf;

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl run -it --rm --image=mysql:5.6 --restart=Never mysql-client -- mysql -h mysql -ppassword
If you don't see a command prompt, try pressing enter.

mysql> create database saconf;
Query OK, 1 row affected (0.01 sec)

mysql> █
```

Check the secret.yml file. You might need to change the cluster IP address which you connect to.

Then deploy the application and check that the connection works properly.

```
kubectl apply -f kube/secret.yml
```

```
kubectl apply -f kube/deployment.yml
```

```
kubectl apply -f kube/service.yml
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl apply -f kube/secret.yml
secret "saconf2019-e4-secret" configured
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl apply -f kube/deployment.yml
deployment "app-with-sql-persistence-deployment" configured
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl apply -f kube/service.yml
service "app-with-sql-persistence-service" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# █
```

```
kubectl get pods
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
app-with-sql-persistence-deployment-55f65cbd58-rrkdb 1/1     Running   0           1m
mysql-6b698dfcbb-q8vwm             1/1     Running   0           8m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

Display logs:

kubectl logs <name_of_the_pod>

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl logs app-with-sql-persistence-deployment-55f65cbd58-rrkdb

:: Spring Boot :: (v2.1.1.RELEASE)

2018-12-08 18:43:41.119 INFO 1 --- [main] c.e.s.Saconf2019E4Application : Starting Saconf2019E4Application v0.0.1-SNAPSHOT on app-with-sql-persistence-deployment-55f65cbd58-rrkdb with PID 1 (/app.jar started by root in /)
2018-12-08 18:43:41.356 INFO 1 --- [main] c.e.s.Saconf2019E4Application : No active profile set, falling back to default profiles: default
2018-12-08 18:44:04.555 INFO 1 --- [main] s.d.r.c.RepositoryConfigurationDelegate : Bootstrapping Spring Data repositories in DEFAULT mode.
2018-12-08 18:44:05.510 INFO 1 --- [main] s.d.r.c.RepositoryConfigurationDelegate : Finished Spring Data repository scanning in 901ms. Found 1 repository interfaces.
2018-12-08 18:44:19.291 INFO 1 --- [main] trationDelegate$BeanPostProcessorChecker : Bean 'org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration' of type [org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration$$EnhancerBySpringCGLIB$$c1025e26] is not eligible for getting processed by all BeanPostProcessors (for example: not eligible for auto-proxying)
2018-12-08 18:44:26.743 INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2018-12-08 18:44:27.390 INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2018-12-08 18:44:27.391 INFO 1 --- [main] org.apache.catalina.core.StandardEngine : Starting Servlet Engine: Apache Tomcat/9.0.13
2018-12-08 18:44:27.591 INFO 1 --- [main] o.a.catalina.core.AprLifecycleListener : The APR based Apache Tomcat Native library which allows optimal performance in production environments was not found on the java.library.path: [/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64/server:/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64:/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64:/usr/java/packages/lib/amd64:/usr/lib64:/lib:/usr/lib]
2018-12-08 18:44:30.094 INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring embedded WebApplicationContext
2018-12-08 18:44:30.095 INFO 1 --- [main] o.s.web.context.ContextLoader : Root WebApplicationContext: initialization completed in 47934 ms
2018-12-08 18:44:34.829 INFO 1 --- [main] o.f.c.internal.license.VersionPrinter : Flyway Community Edition 5.2.3 by Boxfuse
2018-12-08 18:44:34.902 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Starting...
2018-12-08 18:44:39.746 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Start completed.
2018-12-08 18:44:39.819 INFO 1 --- [main] o.f.c.internal.database.DatabaseFactory : Database: jdbc:mysql://172.31.40.33:30001/saconf (MySQL 5.6)
2018-12-08 18:44:43.346 INFO 1 --- [main] o.f.core.internal.command.DbValidate : Successfully validated 1 migration (execution time 00:01:721s)
2018-12-08 18:44:44.104 INFO 1 --- [main] o.f.c.i.s.JdbcTableSchemaHistory : Creating Schema History table: 'saconf`.`flyway_schema_history'
2018-12-08 18:44:44.365 INFO 1 --- [main] o.f.core.internal.command.DbMigrate : Current version of schema 'saconf': << Empty Schema >>
2018-12-08 18:44:44.383 INFO 1 --- [main] o.f.core.internal.command.DbMigrate : Migrating schema 'saconf' to version 1 - CREATE USERS TABLE
2018-12-08 18:44:44.749 INFO 1 --- [main] o.f.core.internal.command.DbMigrate : Successfully applied 1 migration to schema 'saconf' (execution time 00:00.645s)
2018-12-08 18:44:46.685 INFO 1 --- [main] o.hibernate.jpa.internal.util.LogHelper : HHH000204: Processing PersistenceUnitInfo [
    name: default
    ...]
2018-12-08 18:44:47.431 INFO 1 --- [main] org.hibernate.Version : HHH0000412: Hibernate Core (5.3.7.Final)
2018-12-08 18:44:47.433 INFO 1 --- [main] org.hibernate.cfg.Environment : HHH000206: hibernate.properties not found
2018-12-08 18:44:51.766 INFO 1 --- [main] o.hibernte.annotations.common.Version : HCN0000001: Hibernate Commons Annotations (5.0.4.Final)
2018-12-08 18:44:54.049 INFO 1 --- [main] org.hibernate.dialect.Dialect : HH0000400: Using dialect: org.hibernate.dialect.MySQL5Dialect
2018-12-08 18:45:00.914 INFO 1 --- [main] j.LocalContainerEntityManagerFactoryBean : Initialized JPA EntityManagerFactory for persistence unit 'default'
2018-12-08 18:45:08.806 INFO 1 --- [main] o.s.s.concurrent.ThreadPoolTaskExecutor : Initializing ExecutorService 'applicationTaskExecutor'
2018-12-08 18:45:09.027 WARN 1 --- [main] aWebConfiguration$JpaWebMvcConfiguration : spring.jpa.open-in-view is enabled by default. Therefore, database queries may be performed during view rendering. Explicitly configure spring.jpa.open-in-view to disable this warning
2018-12-08 18:45:09.957 INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port(s): 8080 (http) with context path ''
2018-12-08 18:45:09.966 INFO 1 --- [main] c.e.s.Saconf2019E4Application : Started Saconf2019E4Application in 103.13 seconds (JVM running for 115.064)
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

kubectl run -it --rm --image=mysql:5.6 --restart=Never mysql-client -- mysql -h mysql -ppassword

use saconf;

show tables;

describe user;


```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl run -it --rm --image=mysql:5.6 --restart=Never mysql-client -- mysql -h mysql -ppassword

If you don't see a command prompt, try pressing enter.

mysql> use saconf;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_saconf |
+-----+
| flyway_schema_history |
| user |
+-----+
2 rows in set (0.00 sec)

mysql> describe user;
+-----+-----+-----+-----+-----+-----+
| Field | Type | Null | Key | Default | Extra |
+-----+-----+-----+-----+-----+-----+
| id | bigint(20) | NO | PRI | NULL | auto_increment |
| username | varchar(100) | NO | UNI | NULL |
+-----+-----+-----+-----+-----+-----+
2 rows in set (0.03 sec)

mysql>

```

Use vi to change the version number in the deployment file to upgrade to version v2.

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# vi kube/deployment.yml
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# cat kube/deployment.yml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-with-sql-persistence-deployment
  labels:
    app: app-with-sql-persistence-deployment
spec:
  replicas: 1
  selector:
    matchLabels:
      app: app-with-sql-persistence-deployment
  template:
    metadata:
      labels:
        app: app-with-sql-persistence-deployment
    spec:
      containers:
        - name: app-with-sql-persistence-deployment
          image: laurentiuspilca/saconf2019-e4:v2
          env:
            - name: spring.datasource.url
              valueFrom:
                secretKeyRef:
                  name: saconf2019-e4-secret
                  key: spring.datasource.url
            - name: spring.datasource.username
              valueFrom:
                secretKeyRef:
                  name: saconf2019-e4-secret
                  key: spring.datasource.username
            - name: spring.datasource.password
              valueFrom:
                secretKeyRef:
                  name: saconf2019-e4-secret
                  key: spring.datasource.password
          ports:
            - containerPort: 8080
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#

```

Apply again the deployment:

kubectl apply -f kube/deployment.yml

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl apply -f kube/deployment.yml
deployment "app-with-sql-persistence-deployment" configured
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

Check the pods:

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
app-with-sql-persistence-deployment-568c469ffc-rvz7j  1/1     Running   0           27s
mysql-6b698dfcbb-g8vwm              1/1     Running   0           16m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl logs app-with-sql-persistence-deployment-568c469ffc-rvz7j

:: Spring Boot ::
(v2.1.1.RELEASE)

2018-12-08 18:53:00.296 INFO 1 --- [main] c.e.s.Saconf2019E4Application : Starting Saconf2019E4Application v0.0.1-SNAPSHOT on app-with-sql-persistence-deployment-568c469ffc-rvz7j with PID 1 (/app.jar started by root in /)
2018-12-08 18:53:00.426 INFO 1 --- [main] c.e.s.Saconf2019E4Application : No active profile set, falling back to default profiles: default
2018-12-08 18:53:24.624 INFO 1 --- [main] .s.d.r.c.RepositoryConfigurationDelegate : Bootstrapping Spring Data repositories in DEFAULT mode.
2018-12-08 18:53:26.261 INFO 1 --- [main] .s.d.r.c.RepositoryConfigurationDelegate : Finished Spring Data repository scanning in 1442ms. Found 1 repository interfaces.
2018-12-08 18:53:35.325 INFO 1 --- [main] trationDelegate$BeanPostProcessorChecker : Bean 'org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration' of type [org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration$$EnhancerBySpringCGLib$$ac780933] is not eligible for getting processed by all BeanPostProcessors (for example: not eligible for auto-proxying)
2018-12-08 18:53:41.679 INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2018-12-08 18:53:42.055 INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2018-12-08 18:53:42.055 INFO 1 --- [main] org.apache.catalina.core.StandardEngine : Starting Servlet Engine: Apache Tomcat/9.0.13
2018-12-08 18:53:42.232 INFO 1 --- [main] o.a.catalina.core.AprLifecycleListener : The APR based Apache Tomcat Native library which allows optimal performance in production environments was not found on the java.library.path: [/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64/server:/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64/lib64:/lib:/usr/lib]
2018-12-08 18:53:45.596 INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring embedded WebApplicationContext
2018-12-08 18:53:45.597 INFO 1 --- [main] o.s.web.context.ContextLoader : Root WebApplicationContext: initialization completed in 43396 ms
2018-12-08 18:53:48.249 INFO 1 --- [main] o.f.c.internal.license.VersionPrinter : Flyway Community Edition 5.2.3 by Boxfuse
2018-12-08 18:53:48.989 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Starting...
2018-12-08 18:53:55.113 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Start completed.
2018-12-08 18:53:55.121 INFO 1 --- [main] o.f.c.internal.database.DatabaseFactory : Database: jdbc:mysql://172.31.40.33:30001/saconf (MySQL 5.6)
2018-12-08 18:53:58.667 INFO 1 --- [main] o.f.c.internal.command.DbValidate : Successfully validated 2 migrations (execution time 00:00.612s)
2018-12-08 18:53:58.760 INFO 1 --- [main] o.f.c.internal.command.DbMigrate : Current version of schema 'saconf': 1
2018-12-08 18:53:58.775 INFO 1 --- [main] o.f.c.internal.command.DbMigrate : Migrating schema 'saconf' to version 2 - ADD PASSWORD COLUMN
2018-12-08 18:53:59.839 INFO 1 --- [main] o.f.c.internal.command.DbMigrate : Successfully applied 1 migration to schema 'saconf' (execution time 00:01.137s)
2018-12-08 18:54:01.416 INFO 1 --- [main] o.hibernate.jpa.internal.util.LogHelper : HHH000204: Processing PersistenceUnitInfo [
    name: default
    ...]
2018-12-08 18:54:02.470 INFO 1 --- [main] org.hibernate.Version : HHH0000412: Hibernate Core (5.3.7.Final)
2018-12-08 18:54:02.492 INFO 1 --- [main] org.hibernate.cfg.Environment : HHH0000206: hibernate.properties not found
2018-12-08 18:54:04.035 INFO 1 --- [main] o.hibernate.annotations.common.Version : HCANN0000001: Hibernate Commons Annotations (5.0.4.Final)
2018-12-08 18:54:11.361 INFO 1 --- [main] org.hibernate.dialect.Dialect : HHH0000400: Using dialect: org.hibernate.dialect.MySQL5Dialect
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

kubectl rollout undo deployment app-with-sql-persistence-deployment

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl rollout undo deployment app-with-sql-persistence-deployment
deployment "app-with-sql-persistence-deployment" rolled back
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4#
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e4
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e4# kubectl logs app-with-sql-persistence-deployment-55f65cbd58-qt4vj
```

[illegible]

```

2018-12-08 18:57:04.879 INFO 1 --- [in] c.e.s.Sacon2019E4Application : Starting Sacon2019E4Application v0.0.1-SNAPSHOT on app-with-sql-persistence-deployment-55f6
Scd58-gt4v] with PID 1 (/app.jar started by root in f)
2018-12-08 18:57:05.041 INFO 1 --- [main] c.e.s.Sacon2019E4Application : No active profile set, falling back to default profiles: default
2018-12-08 18:57:27.150 INFO 1 --- [main] s.d.r.c.RepositoryConfigurationDelegate : Bootstrapping Spring Data repositories in DEFAULT mode.
2018-12-08 18:57:29.145 INFO 1 --- [main] s.d.r.c.RepositoryConfigurationDelegate : Finished Spring Data repository scanning in 965ms. Found 1 repository interfaces.
2018-12-08 18:57:36.559 INFO 1 --- [main] trationDelegateBeanPostProcessorChecker : Bean 'org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration' of type [org.springframework.transaction.annotation.ProxyTransactionManagementConfiguration$Enhancer$BySpringCGIIBSS$fc704a5] is not eligible for getting processed by all BeanPostProcessors (for example: not eligible for auto-proxying)
2018-12-08 18:57:49.698 INFO 1 --- [main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port(s): 8080 (http)
2018-12-08 18:57:50.389 INFO 1 --- [main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2018-12-08 18:57:50.445 INFO 1 --- [main] org.apache.catalina.core.StandardEngine : Starting Servlet Engine: Apache Tomcat/9.0.13
2018-12-08 18:57:52.365 INFO 1 --- [main] o.a.catalina.core.AprLifecycleListener : The APR based Apache Tomcat Native library which allows optimal performance in production environments was not found on the java.library.path: [/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64/server:/usr/lib/jvm/java-1.8-openjdk/jre/lib/amd64:/usr/lib/jvm/java-1.8-openjdk/jre/..:/lib/m64:/usr/java/packages/lib/amd64:/usr/lib64:/lib64:/lib:/usr/lib]
2018-12-08 18:57:55.094 INFO 1 --- [main] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring embedded WebApplicationContext
2018-12-08 18:57:55.096 INFO 1 --- [main] org.web.context.ContextLoader : Root WebApplicationContext: initialization completed in 4898 ms
2018-12-08 18:57:55.666 INFO 1 --- [main] o.f.c.internal.license.VersionPrinter : Flyway Community Edition 5.2.2 by BoXuse
2018-12-08 18:57:59.708 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Starting...
2018-12-08 18:58:05.574 INFO 1 --- [main] com.zaxxer.hikari.HikariDataSource : HikariPool-1 - Start completed.
2018-12-08 18:58:05.651 INFO 1 --- [main] o.f.c.internal.database.DatabaseFactory : Database: jdbc:ucsql://172.31.40.33:30001/saconf (MySQL 5.6)
2018-12-08 18:58:08.556 INFO 1 --- [main] o.f.c.internal.command.DBValidate : Successfully validated 2 migrations (execution time 00:00.556s)
2018-12-08 18:58:09.651 INFO 1 --- [main] o.f.c.internal.command.DBMigrate : Current version of schema 'saconf' is 2
2018-12-08 18:58:09.651 INFO 1 --- [main] o.f.c.internal.command.DBMigrate : Schema 'saconf' has a version (2) that is newer than the latest available migration (1) !
2018-12-08 18:58:09.651 INFO 1 --- [main] o.f.c.internal.command.DBMigrate : Schema 'saconf' is up to date. No migration necessary.
2018-12-08 18:58:11.180 INFO 1 --- [main] o.hibernate.jpa.internal.util.LogHelper : HH000204: Processing PersistenceUnitInfo [
name: default
...]
2018-12-08 18:58:12.405 INFO 1 --- [main] org.hibernate.cfg : HH0000412: Hibernate Core [5.2.7.Final]
2018-12-08 18:58:12.621 INFO 1 --- [main] org.hibernate.cfg.Environment : HH0000206: Hibernate properties not found
2018-12-08 18:58:16.678 INFO 1 --- [main] o.hibernate.annotations.common.Version : HCANN000001: Hibernate Commons Annotations [5.0.4.Final]
2018-12-08 18:58:20.681 INFO 1 --- [main] org.hibernate.dialect.Dialect : HH0000400: Using dialect: org.hibernate.dialect.MySQL5Dialect
2018-12-08 18:58:27.418 INFO 1 --- [main] j.LocalContainerEntityManagerFactoryBean : Initialized JPA EntityManagerFactory for persistence unit 'default'
root@bz-172-31-40-33:/var/saconf2019/saconf2019-e4#

```

6. The actuator – check the healthiness of your container

In the early days of software, the application knew about the environment. Now the relationship is bidirectional. The environment must also know about the application. Frameworks designed for microservices architectures - like Spring Boot in java - come with implementations like the actuator to facilitate this.

```
kubectl apply -f kube/deployment.yml
```

```
kubectl apply -f kube/service.yml
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5# kubectl apply -f kube/deployment.yml
deployment "health-check-app-deployment" created
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5# kubectl apply -f kube/service.yml
service "health-check-app-service" created
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5#
```

```
kubectl get services
```

```
curl http://localhost:<port>/actuator/health
```

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5# kubectl get services
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)                                AGE
health-check-app-service            LoadBalancer        10.105.60.204    <pending>         8080:31296/TCP,8081:31785/TCP          1m
kubernetes                           ClusterIP             10.96.0.1         <none>            443/TCP                                13d
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5# curl http://localhost:31785/actuator/health
{"status": "UP"}root@ip-172-31-40-33: /var/saconf2019/saconf2019-e5#
```

```
kubectl apply -f kube/mysql-deployment.yml
```

```
kubectl apply -f kube/mysql-service.yml
```

```
kubectl apply -f kube/secret.yml
```

```
kubectl apply -f kube/deployment.yml
```

```
kubectl apply -f kube/service.yml
```

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e6
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl apply -f kube/mysql-deployment.yml
deployment "mysql" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl apply -f kube/mysql-service.yml
service "mysql" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl apply -f kube/secret.yml
secret "saconf2019-e6-secret" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl apply -f kube/deployment.yml
deployment "health-check-app-db-deployment" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl apply -f kube/service.yml
service "health-check-app-db-service" created
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# █

```

kubectl get pods

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e6
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
health-check-app-db-deployment-57c585bbc9-hzsmx  1/1     Running   0           1m
mysql-6b698dfcbb-zfxd5                1/1     Running   0           2m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# █

```

curl http://localhost:<port>/actuator/health

```

root@ip-172-31-40-33: /var/saconf2019/saconf2019-e6
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl get services
NAME                                TYPE           CLUSTER-IP      EXTERNAL-IP      PORT(S)                                     AGE
health-check-app-db-service        LoadBalancer  10.101.157.64   <pending>        8080:30792/TCP,8081:31570/TCP             2m
kubernetes                         ClusterIP      10.96.0.1       <none>           443/TCP                                    13d
mysql                              LoadBalancer  10.98.29.66     <pending>        3306:30001/TCP                             3m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# curl http://localhost:31570/actuator/health
{"status":"UP","details":{"db":{"status":"UP","details":{"database":"MySQL","hello":1}},"diskSpace":{"status":"UP",
ot@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# █

```

kubectl delete deployments mysql

kubectl delete services mysql

curl http://localhost:<port>/actuator/health

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e6
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# curl http://localhost:31570/actuator/health
{"status":"DOWN","details":{"db":{"status":"DOWN","details":{"error":"org.springframework.jdbc.CannotGetJdbc
ql.SQLTransientConnectionException: HikariPool-1 - Connection is not available, request timed out after 3000
4,\"threshold\":10485760}}}}root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6#
```

kubectl get pods

```
root@ip-172-31-40-33: /var/saconf2019/saconf2019-e6
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
health-check-app-db-deployment-57c585bbc9-hzsmx  0/1     Running   1           9m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e6#
```

7. An unfortunate choice - having multiple containers in the same pod

We usually think of a pod as having only one container. But a pod can actually have multiple containers. Is it a good practice to deploy related containers in the same pod?

Let's first start with seeing how is this possible.

```
kubectl apply -f kube/deployment.yml
```

```
kubectl apply -f kube/service.yml
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7# kubectl apply -f kube/deployment.yml
deployment "multiple-containers-deployment" unchanged
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7# kubectl apply -f kube/service.yml
service "multiple-containers-app" unchanged
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7#
```

```
kubectl get services
```

```
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7# kubectl get services
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes           ClusterIP     10.96.0.1     <none>         443/TCP          14d
multiple-containers-app LoadBalancer  10.104.42.124 <pending>      8080:30618/TCP,9090:31413/TCP 4m
root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7# curl http://localhost:30618/container1
From container 1root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7# curl http://localhost:31413/container2
From container 2root@ip-172-31-40-33:/var/saconf2019/saconf2019-e7#
```

8. Limits and auto-scaling

Limitation is important. We always have to define the request and limit resources for our containers. This can be done in the deployment.yml file as shows in the example 8.

Apply the yml files from the kube folder to run the example. As you have observed we also have here the definition of an auto-scaler within the file hpa.yml

Once we have defined limits we can use the metrics server to display the current status.

kubectl top pods

```
root@ip-172-31-43-94: /var/saconf2019/saconf2019-e8
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# kubectl top pods
NAME                                                    CPU(cores)   MEMORY(bytes)
autoscaling-app-deployment-659f89f4c9-46mnh           0m           103Mi
autoscaling-app-deployment-659f89f4c9-fpc7z           0m           104Mi
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8#
```

```
root@ip-172-31-43-94: /var/saconf2019/saconf2019-e8
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# kubectl get hpa
NAME                REFERENCE                                TARGETS   MINPODS   MAXPODS   REPLICAS   AGE
ex8-pod-autoscaler  Deployment/autoscaling-app-deployment  0%/5%     2          10         2          24s
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8#
```

We can use Apache Benchmark to stress a little our service and see the horizontal pod auto-scaler in action.

apt-get install apache2-utils


```
root@ip-172-31-43-94: /var/saconf2019/saconf2019-e8
```

```
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# kubectl get services
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
autoscaling-app     LoadBalancer       10.104.225.210  <pending>        8080:31825/TCP   10m
kubernetes           ClusterIP           10.96.0.1       <none>           443/TCP          11m
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# curl http://localhost:31825/test
okroot@ip-172-31-43-94:/var/saconf2019/saconf2019-e8#
```

```
root@ip-172-31-43-94: /var/saconf2019/saconf2019-e8
```

```
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# ab -n 10 -c 1 http://localhost:31825/test
This is ApacheBench, Version 2.3 <$Revision: 1706008 $>
Copyright 1996 Adam Twiss, Zeus Technology Ltd, http://www.zeustech.net/
Licensed to The Apache Software Foundation, http://www.apache.org/
```

```
Benchmarking localhost (be patient).....done
```

```
Server Software:
Server Hostname: localhost
Server Port: 31825

Document Path: /test
Document Length: 2 bytes

Concurrency Level: 1
Time taken for tests: 65.832 seconds
Complete requests: 10
Failed requests: 0
Total transferred: 1340 bytes
HTML transferred: 20 bytes
Requests per second: 0.15 [#/sec] (mean)
Time per request: 6583.210 [ms] (mean)
Time per request: 6583.210 [ms] (mean, across all concurrent requests)
Transfer rate: 0.02 [Kbytes/sec] received
```

```
Connection Times (ms)
      min  mean[+/-sd] median   max
Connect:    0      0  0.0/0.0      0      0
Processing: 2411  6583 3509.8   8333  12194
Waiting:    2411  6563 3495.5   8240  12190
Total:      2411  6583 3509.8   8333  12194
```

```
Percentage of the requests served within a certain time (ms)
```

```
50%   8333
66%   9005
75%   9473
80%   9602
90%  12194
95%  12194
98%  12194
99%  12194
100% 12194 (longest request)
```

```
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# kubectl get hpa
NAME                REFERENCE                TARGETS  MINPODS  MAXPODS  REPLICAS  AGE
ex8-pod-autoscaler  Deployment/autoscaling-app-deployment  35%/5%   2        10       10        5m
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8#
```

 root@ip-172-31-43-94: /var/saconf2019/saconf2019-e8

```
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8# kubectl get pods
NAME                                READY    STATUS    RESTARTS   AGE
autoscaling-app-deployment-659f89f4c9-46mnh  1/1      Running   0          14m
autoscaling-app-deployment-659f89f4c9-4rk78  1/1      Running   0          2m
autoscaling-app-deployment-659f89f4c9-9st25  1/1      Running   0          3m
autoscaling-app-deployment-659f89f4c9-c2z9l  1/1      Running   0          2m
autoscaling-app-deployment-659f89f4c9-cst42  1/1      Running   0          3m
autoscaling-app-deployment-659f89f4c9-fpc7z  1/1      Running   0          14m
autoscaling-app-deployment-659f89f4c9-lgc9r  1/1      Running   0          3m
autoscaling-app-deployment-659f89f4c9-pp4rk  1/1      Running   0          3m
autoscaling-app-deployment-659f89f4c9-zfn76  1/1      Running   0          3m
autoscaling-app-deployment-659f89f4c9-zg4gf  1/1      Running   0          3m
root@ip-172-31-43-94:/var/saconf2019/saconf2019-e8#
```